

Lab 3

Course: Operating Systems

Toan V. Tran, Pham Trung Kien, Phuong-Duy Nguyen
Email: viettoantran98@gmail.com

April 5, 2021

Goal: This lab helps student to understand the definition of process, and how an operating system can manage the execution of a process.

Content In detail, this lab reflects the theory of process into practical exercises. For example:

- How to retrieve the information of running process? How is PCB table controlled by OS?
- Create a program with multiple processes.

Result After doing this lab, student can distinguish a program and a process. They can create a program with multiple process and retrieve the information of processes.

Requirement Student need to review the theory of process in operating system.

1. INTRODUCTION

Informally, as mentioned earlier, a process is a program in execution. A process is more than the program code, which is sometimes known as the text section. It also includes the current activity, as represented by the value of the program counter and

the contents of the processor's registers. A process generally also includes the process stack, which contains temporary data (such as function parameters, return addresses, and local variables), and a data section, which contains global variables. A process may also include a heap, which is memory that is dynamically allocated during process run time.

As a process executes, it changes state. The state of a process is defined in part by the current activity of that process.

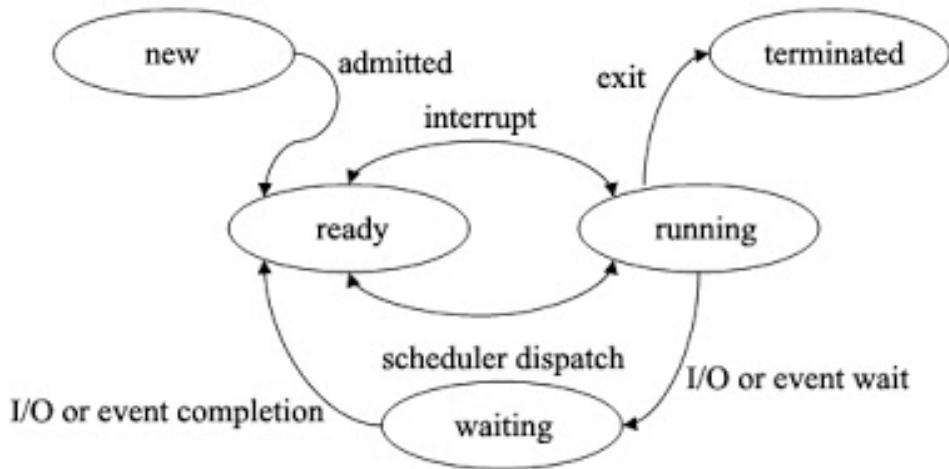


Figure 1.1: Diagram of process state.

To keep track of processes, the operating system maintains a process table (or list). Each entry in the process table corresponds to a particular process and contains fields with information that the kernel needs to know about the process. This entry is called a Process Control Block (PCB). Some of these fields on a typical Linux/Unix system PCB are:

- Machine state (registers, program counter, stack pointer)
- Parent process and a list of child processes
- Process state (ready, running, blocked)
- Event descriptor if the process is blocked
- Memory map (where the process is in memory)
- Open file descriptors
- Owner (user identifier). This determines access privileges & signaling privileges

- Scheduling parameters
- Signals that have not yet been handled
- Timers for accounting (time & resource utilization)
- Process group (multiple processes can belong to a common group)

A process is identified by a unique number called the process ID (PID). Some operating systems (notably UNIX-derived systems) have a notion of a process group. A process group is just a way to lump related processes together so that related processes can be signaled. Every process is a member of some process group.

2. HOW DO WE FIND THE PROCESS'ID AND GROUP?

A process can find its process ID with the *getpid* system call. It can find its process group number with the *getpgrp* system call, and it can find its parent's process ID with *getppid*. For example:

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4
5 int main(int argc, char **argv) {
6     printf("Process_ID:%d\n", getpid());
7     printf("Parent_process_ID:%d\n", getppid());
8     printf("My_group:%d\n", getpgrp());
9
10    return 0;
11 }
```

3. CREATING A PROCESS

The *fork* system call clones a process into two processes running the same code. *Fork* returns a value of 0 to the child and a value of the process ID number (pid) to the parent. A value of -1 is returned on failure.

```

1 #include <stdio.h>
2 #include <stdlib.h>
3 #include <unistd.h>
4
5 int main(int argc, char **argv) {
6     switch (fork()) {
7         case 0:
8             printf("I_am_the_child:_pid=%d\n", getpid());
```

```

9         break;
10    default:
11        printf("I am the parent: pid=%d\n", getpid());
12        break;
13    case -1:
14        perror("Fork failed");
15    }
16    return 0;
17 }
```

4. BASICS OF MULTI-PROCESS PROGRAMMING

MULTI-PROCESS PROGRAMMING Implement a program using `fork()` command to create child process. Student can check the number of forked processes using `ps` command.

```

1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <unistd.h>      /* defines fork(), and pid_t. */
4
5 int main(int argc, char ** argv) {
6
7     pid_t child_pid;
8
9     /* lets fork off a child process... */
10    child_pid = fork();
11
12    /* check what the fork() call actually did */
13    if (child_pid == -1) {
14        perror("fork"); /* print a system-defined error message */
15        exit(1);
16    }
17
18    if (child_pid == 0) {
19        /* fork() succeeded, we're inside the child process */
20        printf("Hello,\n");
21        fflush(stdout);
22    }
23    else {
24        /* fork() succeeded, we're inside the parent process */
25        printf("World!\n");
26        fflush(stdout);
27    }
28 }
```

```
29     return 0;  
30 }
```

COMPILE AND RUN THE PROGRAM observing the output of the program and giving the conclusion about the order of letters “Hello, World!”.

This is because of the independent among processes, so that the order of execution is not guaranteed. It need a mechanism to suspend the process until its child process finished. The system call `wait()` is used in this case. The example of `wait()` is presented in appendix A.

In case of expanding the result to implement the reverse order which child process is suspended until its father process is finished. This case is not recommended due to the **orphan process** status. An example of such programs is presented in appendix B

5. RETRIEVE THE CODE SEGMENT OF PROCESS

5.1. THE INFORMATION OF PROCESS

You need a program with the execution time being enough long.

```
1 /* Source code of loop_process.c */  
2 #include <stdlib.h>  
3 #include <stdio.h>  
4  
5 int main(int argc, char ** argv) {  
6     int timestamp=0;  
7     while(1){  
8         printf("Time: %5d\n", timestamp++);  
9         sleep(1);  
10    }  
11    return 0;  
12 }
```

```
$ gcc -o loop_process loop_process.c  
$ ./loop_process  
Time:      0  
Time:      1  
Time:      2  
Time:      3  
Time:      4  
...
```

FIND PROCESS ID **ps** command is used to find (process ID - pid) of a process.

```
$ ps -a | grep loop_process
7277 tc      ./loop_process
7279 tc      grep loop_process
```

RETRIEVE PCB INFORMATION OF PROCESS on Linux system, each running process is reflected in the folder of filesystem **/proc**. Information of each process is associated with the folder called **/proc/<pid>**, where pid is process ID. For example, with the *pid* of the process above, **./loop_process** the directory containing the information of this process is **/proc/7277**.

```
$ ls /proc/<pid>
autogroup      environ          mountstats    smaps
auxv          exe              net/           stat
cgroup         fd/             ns/            statm
clear_refs     fdinfo/        oom_adj       status
cmdline        limits          oom_score    syscall
comm           maps            oom_score_adj task/
coredump_filter mem            pagemap
cpuset          mountinfo      personality
cwd            mounts          root
```

Using commands such as **cat**, **vi** to show the information of processes.

```
$ cat /proc/<pid>/cmdline
./loop_process

$ cat /proc/<pid>/status
Name:  loop_process
State: S (sleeping)
Tgid: 7277
Pid: 7277
PPid: 6760
TracerPid: 0
Uid: 1001 1001 1001 1001
Gid: 50 50 50 50
FDSize: 32
...
```

5.2. COMPARING THE CODE SEGMENT OF PROCESS AND PROGRAM
/PROC/<PID>/maps stores the information of code in **maps**

```
$ cat /proc/<pid>/maps
```

```

08048000-08049000 r-xp 00000000 00:01 30895 /home/.../loop_process
08049000-0804a000 rwxp 00000000 00:01 30895 /home/.../loop_process
b75e0000-b75e1000 rwxp 00000000 00:00 0
b75e1000-b76f8000 r-xp 00000000 00:01 646 /lib/libc-2.17.so
b76f8000-b76fa000 r-xp 00116000 00:01 646 /lib/libc-2.17.so
b76fa000-b76fb000 rwxp 00118000 00:01 646 /lib/libc-2.17.so
b76fb000-b76fe000 rwxp 00000000 00:00 0
b7705000-b7707000 rwxp 00000000 00:00 0
b7707000-b7708000 r-xp 00000000 00:00 0 [vdso]
b7708000-b7720000 r-xp 00000000 00:01 648 /lib/ld-2.17.so
b7720000-b7721000 r-xp 00017000 00:01 648 /lib/ld-2.17.so
b7721000-b7722000 rwxp 00018000 00:01 648 /lib/ld-2.17.so
bf9a8000-bf9c9000 rw-p 00000000 00:00 0 [stack]

```

COMPARING WITH THE PROGRAM (executable binary file) using *ldd* to read executable binary file and *readelf* to list libraries that are used.

```

$ ldd loop_process
    linux-gate.so.1 (0xb77dc000)
    libc.so.6 => /lib/libc.so.6 (0xb76b6000)
    /lib/ld-linux.so.2 (0xb77dd000)

$ readelf -Ws /lib/libc.so.6 | grep sleep
388: 000857f0 105 FUNC WEAK  DEFAULT 11 nanosleep@@GLIBC_2.0
660: 000857f0 105 FUNC WEAK  DEFAULT 11 __nanosleep@@GLIBC_2.2.6
803: 000bda48 121 FUNC GLOBAL DEFAULT 11 clock_nanosleep@@GLIBC_2.17
1577: 000bda48 121 FUNC GLOBAL DEFAULT 11 __clock_nanosleep@@GLIBC_PRIVATE
1651: 000aa8f8  43 FUNC GLOBAL DEFAULT 11 usleep@@GLIBC_2.0
1959: 00085564 542 FUNC WEAK  DEFAULT 11 sleep@@GLIBC_2.0

```

Following that, we can see the consistency of code segment between program and process.

6. EXERCISE

6.1. QUESTIONS

- What the output will be at LINE A?

```

#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int value = 5;

```

```

int main()
{
    pid_t pid;
    pid = fork();
    if (pid == 0) { /* child process */
        value += 15;
        return 0;
    }
    else if (pid > 0) { /* parent process */
        wait(NULL);
        printf("PARENT: value=%d", value); /* LINE A */
        return 0;
    }
}

```

2. How many processes are created by the program shown below, including the initial parent process? How many process are created when **n** *fork()* called?

```

#include <stdio.h>
#include <unistd.h>
int main()
{
    /* fork a child process */
    fork();

    /* fork another child process */
    fork();

    /* and fork another */
    fork();

    return 0;
}

```

3. When a process creates a new process using the *fork()* operation, which of the following states is shared between the parent process and the child process? Why?
- Stack
 - Heap
 - Shared memory segments
4. What process id (PID) and process group id are used for?

6.2. PROGRAMMING EXERCISES (REQUIRED)

In Lab 2, you have been familiar with the source code to build a simple shell. In this lab, you still are provided a new source with minor changes. Your tasks are to fill parts in file "fork_shell.c". Students should read comments in "fork_shell.c" to understand variables and functions.

PROBLEM 1 The relationship between processes could be represented by a tree. When a process uses *fork* system call to create another process then the new process is a *child* of this process. This process is the parent of the new process. For examples, if process A uses two *fork* system calls to create two new processes B and C then we could display their relationship by a tree in figure 6.1. B is a child process of A. A is the parent of both B and C.

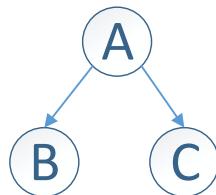


Figure 6.1: Creating processing with *fork()*.

Here is the output of the tree in Figure 6.1:

```
$ ./forkShell
Example Fork Shell>$ tree
My PID is 3773
My PID is 3776, My parent's PID is 3773
...My PID is 3777, My parent's PID is 3773
Example Fork Shell>$ # CTRL + D
Free memory... Quit.. (This line should be printed...!)
```

Write code in "fork_shell.c" to finish the "tree" function that uses *fork* system calls to create processes whose relationship is similar to the one showed in Figure 6.2.

Note: If a process has multiple children then its children must be created from left to right. For example, process A must creates B first then create C and finally D.

After executing the program, we must see the following lines on the screen (Note that the order may be different when running many times.)

PROBLEM 2 (5 POINTS) Hint: Similarly to lab 2, now the variable named "commands" is a pointer to a list of commands. For example: if the input is: "sleep 3; sleep 4; echo abc xyz", you will get:

```
{  
    type == ',';
```

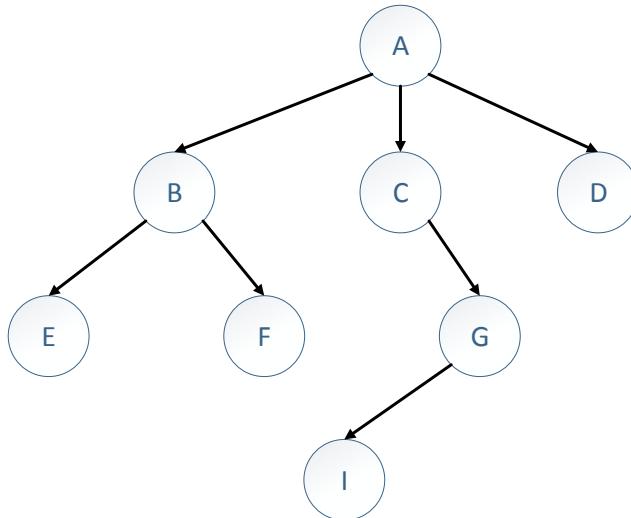


Figure 6.2: The tree of running processes.

```

commands -> [
    [
        ["sleep", "3", NULL],
        ["sleep", "4", NULL],
        ["echo", "abc", "xyz", NULL],
        NULL
    ]
}
  
```

Students should call function `execute_command(commands[i])` to execute a command. Your task is implement three types of command combination:

- SEQUENTIAL_COMM (";") – this command combinations will execute commands sequentially.

```

$ ./forkShell
Shell>$ sleep 3; sleep 4; echo abc
waked after sleeping for 3 seconds
waked after sleeping for 4 seconds
abc
Shell>$ # CTRL + D
Free memory... Quit... (This line should be ...!)
  
```

- BACKGROUND_COMM ("&") – this command combination will execute commands in parallel. The parent process will create many child processes, each child process executes a command. The parent process continues without waiting child processes.

```

$ ./forkShell
Shell>$ sleep 3 & sleep 4 & echo abc
abc xyz
Shell>$ waked after sleeping for 3 seconds
waked after sleeping for 4 seconds
# CTRL + D
Free memory... Quit.. (This line should be ...!)

```

- SEQUENTIAL_BACKGROUND_COMM ("*") – similarly to BACKGROUND_COMM, but the parent will wait until all child processes finish.

```

$ ./forkShell
Shell>$ sleep 3 * sleep 4 * echo abc
abc
waked after sleeping for 3 seconds
waked after sleeping for 4 seconds
Shell>$ # CTRL + D
Free memory... Quit.. (This line should be ...!)

```

SUBMISSION

NOTE: As these exercises are graded automatically, thereby, you need to implement the program by the requirements mentioned above. you must compress all of files (.c, .h, Makefile) into a zip file as:

```

MSSV.zip
├── Question
│   └── answer.txt
└── Program
    ├── Makefile
    ├── fork_shell.c
    ├── parse_command.h
    └── parse_command.c

```

Any cheating code detected will get 0 score. Students are required to show their running program during lab hours.

A. SYSTEM CALL WAIT()

This example uses system call to guarantee the order of running processes.

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <unistd.h>      /* defines fork(), and pid_t. */
4
5 int main(int argc, char ** argv) {
6
7     pid_t child_pid;
8
9     /* lets fork off a child process... */
10    child_pid = fork();
11
12    /* check what the fork() call actually did */
13    if (child_pid == -1) {
14        perror("fork");
15        exit(1);
16    }
17
18    if (child_pid == 0) {
19        /* fork() succeeded, we're inside the child process */
20        printf("Hello ,\n");
21        fflush(stdout);
22    }
23    else {
24        /* fork() succeeded, we're inside the parent process */
25        wait(NULL);           /* wait the child
26        exit */
27        printf("World!\n");
28        fflush(stdout);
29    }
30
31    return 0;
32 }
```

B. SIGNAL

This example illustrate the using of IPC `signal()` and `kill()` routines to suspend child process until its parent process is finished.

```
1 #include <stdlib.h>
2 #include <stdio.h>
3 #include <unistd.h>      /* defines fork(), and pid_t. */
4
5 int main(int argc, char ** argv) {
6
7     pid_t child_pid;
8     sigset_t mask, oldmask;
9
10    /* lets fork off a child process... */
11    child_pid = fork();
12
13    /* check what the fork() call actually did */
14    if (child_pid == -1) {
15        perror("fork"); /* print a system-defined error
16 message */
17        exit(1);
18    }
19
20    if (child_pid == 0) {
21        /* fork() succeeded, we're inside the child process */
22        signal(SIGUSR1, parentdone);           /* set up a signal */
23        /* Set up the mask of signals to temporarily block. */
24        sigemptyset (&mask);
25        sigaddset (&mask, SIGUSR1);
26
27        /* Wait for a signal to arrive. */
28        sigprocmask (SIG_BLOCK, &mask, &oldmask);
29        while (!usr_interrupt)
30            sigsuspend (&oldmask);
31        sigprocmask (SIG_UNBLOCK, &mask, NULL);
32
33
34        printf("World!\n");
35        fflush(stdout);
36    }
37    else {
38        /* fork() succeeded, we're inside the parent process */
39        printf("Hello,~");
```

```
40     fflush( stdout );
41     kill( child_pid , SIGUSR1 );
42     wait( NULL );
43 }
44
45     return 0;
46 }
```

REVISION HISTORY

Revision	Date	Author(s)	Description
1.0	22.09.15	PD Nguyen	created
2.0	19.02.19	TK Pham	Change Exercises
3.0	05.04.21	Toan TV	Change Exercises